



Indian Institute of Technology Jodhpur  
Department of Computer Science and Engineering

## DBS Project

Abhyudaya Tiwari  
B23CS1085

Neeraj Kumar  
B23CS1044

*Supervisor:* Awathare Nitin Niranjan

November 14, 2025

# Contents

<b>1</b>	<b>Project Objectives</b>	<b>1</b>
<b>2</b>	<b>System Architecture</b>	<b>2</b>
2.1	System Layers . . . . .	2
<b>3</b>	<b>File and Code Structure</b>	<b>3</b>
<b>4</b>	<b>Implementation Details: HF (Slotted Page) Layer</b>	<b>4</b>
4.1	Slotted Page Layout . . . . .	4
4.2	Record Insertion . . . . .	4
4.3	Record Deletion . . . . .	4
4.4	Sequential Scanning . . . . .	5
<b>5</b>	<b>Testing and Evaluation</b>	<b>6</b>
5.1	PF Buffer Performance (LRU vs. MRU) . . . . .	6
5.2	HF Space Utilization . . . . .	6
5.3	AM Index Build Performance . . . . .	7
<b>6</b>	<b>Challenges and Learnings</b>	<b>10</b>
6.1	Challenges Faced . . . . .	10
6.2	Take-aways from the Assignment . . . . .	10

# Chapter 1

## Project Objectives

The goal of this project was to implement and evaluate key components of a database storage manager. The system is built in three layers, and the objectives were as follows:

1. **Paged File (PF) Layer:** Implement page buffering for the PF layer with a configurable buffer pool size. The system must support LRU and MRU page replacement strategies, track dirty pages, and report statistics (logical/physical I/O) for different read/write query mixtures.
2. **Heap File (HF) Layer:** Build a slotted-page structure on top of the PF layer to store variable-length records. This layer must support record insertion, deletion, and sequential scanning. Performance metrics, particularly space utilization, must be gathered and compared against static record management.
3. **Access Method (AM) Layer:** Use the AM layer to build a B+-Tree index on a file. The evaluation must compare the performance of building an index on a pre-populated file (using a bulk-loading technique) versus building it incrementally through inserts.

This report focuses primarily on the implementation and evaluation of the **HF (Slotted Page) Layer** and its interaction with the other layers.

# Chapter 2

## System Architecture

The storage manager is designed as a three-layer stack, where each layer provides services to the one above it.

### 2.1 System Layers

- **PF Layer (Paged File):** The lowest layer, providing page-level file I/O. It handles buffer pool management (pinning/unpinning pages) and hides low-level disk operations, presenting a simple page-based storage interface.
- **HF Layer (Heap File):** The middle layer and the focus of this project. It implements a slotted-page organization on top of PF files to efficiently store, retrieve, and manage variable-length records.
- **AM Layer (Access Method):** The highest layer, which builds B+-tree indexes on top of PF files. It uses the HF layer to access the records themselves and provides efficient, indexed access paths.

## Chapter 3

# File and Code Structure

The project is organized into distinct directories representing each layer:

- **pf/**: Contains the Paged File (PF) layer implementation, including the buffer pool manager and page replacement strategies (LRU, MRU).
- **hf/**: Contains the Heap File (HF) layer implementation, which defines the slotted-page structure, record insertion/deletion logic, and scan operations. This was the primary focus of implementation.
- **am/**: Contains the Access Method (AM) layer implementation, providing B+-Tree indexing.
- **test/**: Contains test files used to validate the functionality and measure the performance of each layer (e.g., `hf_test.c`).

# Chapter 4

## Implementation Details: HF (Slotted Page) Layer

The HF layer implements a slotted-page organization for storing variable-length records. Each HF file is composed of pages obtained from the PF layer, and each page is internally structured to manage records efficiently.

### 4.1 Slotted Page Layout

Each page consists of a page header, a slot directory (or slot table), and the record data region.

- **Page Header:** Stores metadata, such as the number of slots in the directory and the offset of the first byte of free space.
- **Slot Directory:** An array of slot entries that grows from the beginning of the page (after the header). Each slot entry contains the **offset** (location within the page) and **length** of a record.
- **Data Region:** Record data is stored contiguously, growing from the end of the page backward.

This "inward-growing" design ensures that free space is always a single contiguous block in the middle of the page.

### 4.2 Record Insertion

When `HF_InsertRecord` is called, the layer searches for a page with enough free space to hold both the new record's bytes and a new slot entry. The process returns a persistent Record Identifier (RID) of `(pageNum, slotNum)`.

### 4.3 Record Deletion

To maintain stable RIDs, records are not physically removed. When `HF_DeleteRecord` is called, the corresponding slot entry is accessed and marked as invalid (e.g., by setting its length or offset




Figure 4.1: Diagram of the Slotted Page Layout.

to -1).

#### 4.4 Sequential Scanning

The HF layer supports sequential scanning via `HF_ScanOpen`, `HF_ScanNext`, and `HF_ScanClose`. The scan iterates through the slot directory of each page and returns the data only from valid (non-deleted) slots.

# Chapter 5

## Testing and Evaluation

The system was validated by testing each layer’s objectives.

### 5.1 PF Buffer Performance (LRU vs. MRU)

The PF layer’s buffer manager was tested with workloads of different read/write query mixtures. The time taken for different operation types (e.g., sequential scan, random access) was measured for both LRU and MRU replacement strategies.




Figure 5.1: PF Layer: Performance (Time vs. Operation Type) for LRU vs. MRU.

### 5.2 HF Space Utilization

The HF layer was tested by loading a dataset of variable-length records. We compared the space efficiency of our slotted-page implementation against a theoretical static, fixed-length

record layout.




Figure 5.2: HF Layer: Space Utilization vs. Assumed Static Record Size.




Figure 5.3: HF Layer: Total Pages Required vs. Assumed Static Record Size.

**Analysis:** As shown in Figures 5.2 and 5.3, the static layout's utilization plummets and page requirement skyrockets as the assumed record size increases, due to massive internal fragmentation. Our slotted-page implementation's performance remains constant and highly efficient, as it only stores the data that is actually present.

### 5.3 AM Index Build Performance

The AM layer's performance was evaluated by comparing three index-building strategies on the Student file, using `roll-no` as the key:

- **Incremental (Random):** Building the index by inserting records in a random order.
- **Build Existing (Incremental Sorted):** Building the index by scanning and inserting records from an already-existing, sorted heap file.

- **Bulk-Loading:** Using a specialized algorithm on a pre-sorted file to build the B+-Tree from the bottom up, which is far more efficient.

We measured build time, total pages used by the index, and subsequent lookup times.



Figure 5.4: AM Layer: Index Construction Time Comparison.




Figure 5.5: AM Layer: Index Pages Used Comparison.

**Analysis:** As shown in the figures, bulk-loading is significantly superior in both build time and space efficiency (fewer pages used), as it creates a more compact tree. All index methods provide near-constant lookup time, which is dramatically faster than a full heap file scan.




Figure 5.6: AM Layer: Index Lookup Time Comparison.  
"inward-growing" design ensures that free space is always a single contiguous block in the middle of the page.




Figure 5.7: AM Layer: Overall comparison

# Chapter 6

## Challenges and Learnings

### 6.1 Challenges Faced

- **Pointer Arithmetic:** The most significant challenge in the HF layer was correctly managing raw byte offsets and pointer arithmetic within the 4KB page buffer.
- **Layer Interaction:** Ensuring that pages were correctly "pinned" (fixed) from the PF layer before being modified by the HF layer, and "unpinned" (unfixed) when done, was essential to prevent data corruption.

### 6.2 Take-aways from the Assignment

- This project provided a deep, practical understanding of how variable-length records are physically stored on disk, demonstrating why slotted pages are a fundamental database structure.
- I gained a clear understanding of the importance of a layered architecture.
- The performance evaluation highlighted the significant space savings of slotted pages and the critical performance benefits of bulk-loading for index creation.